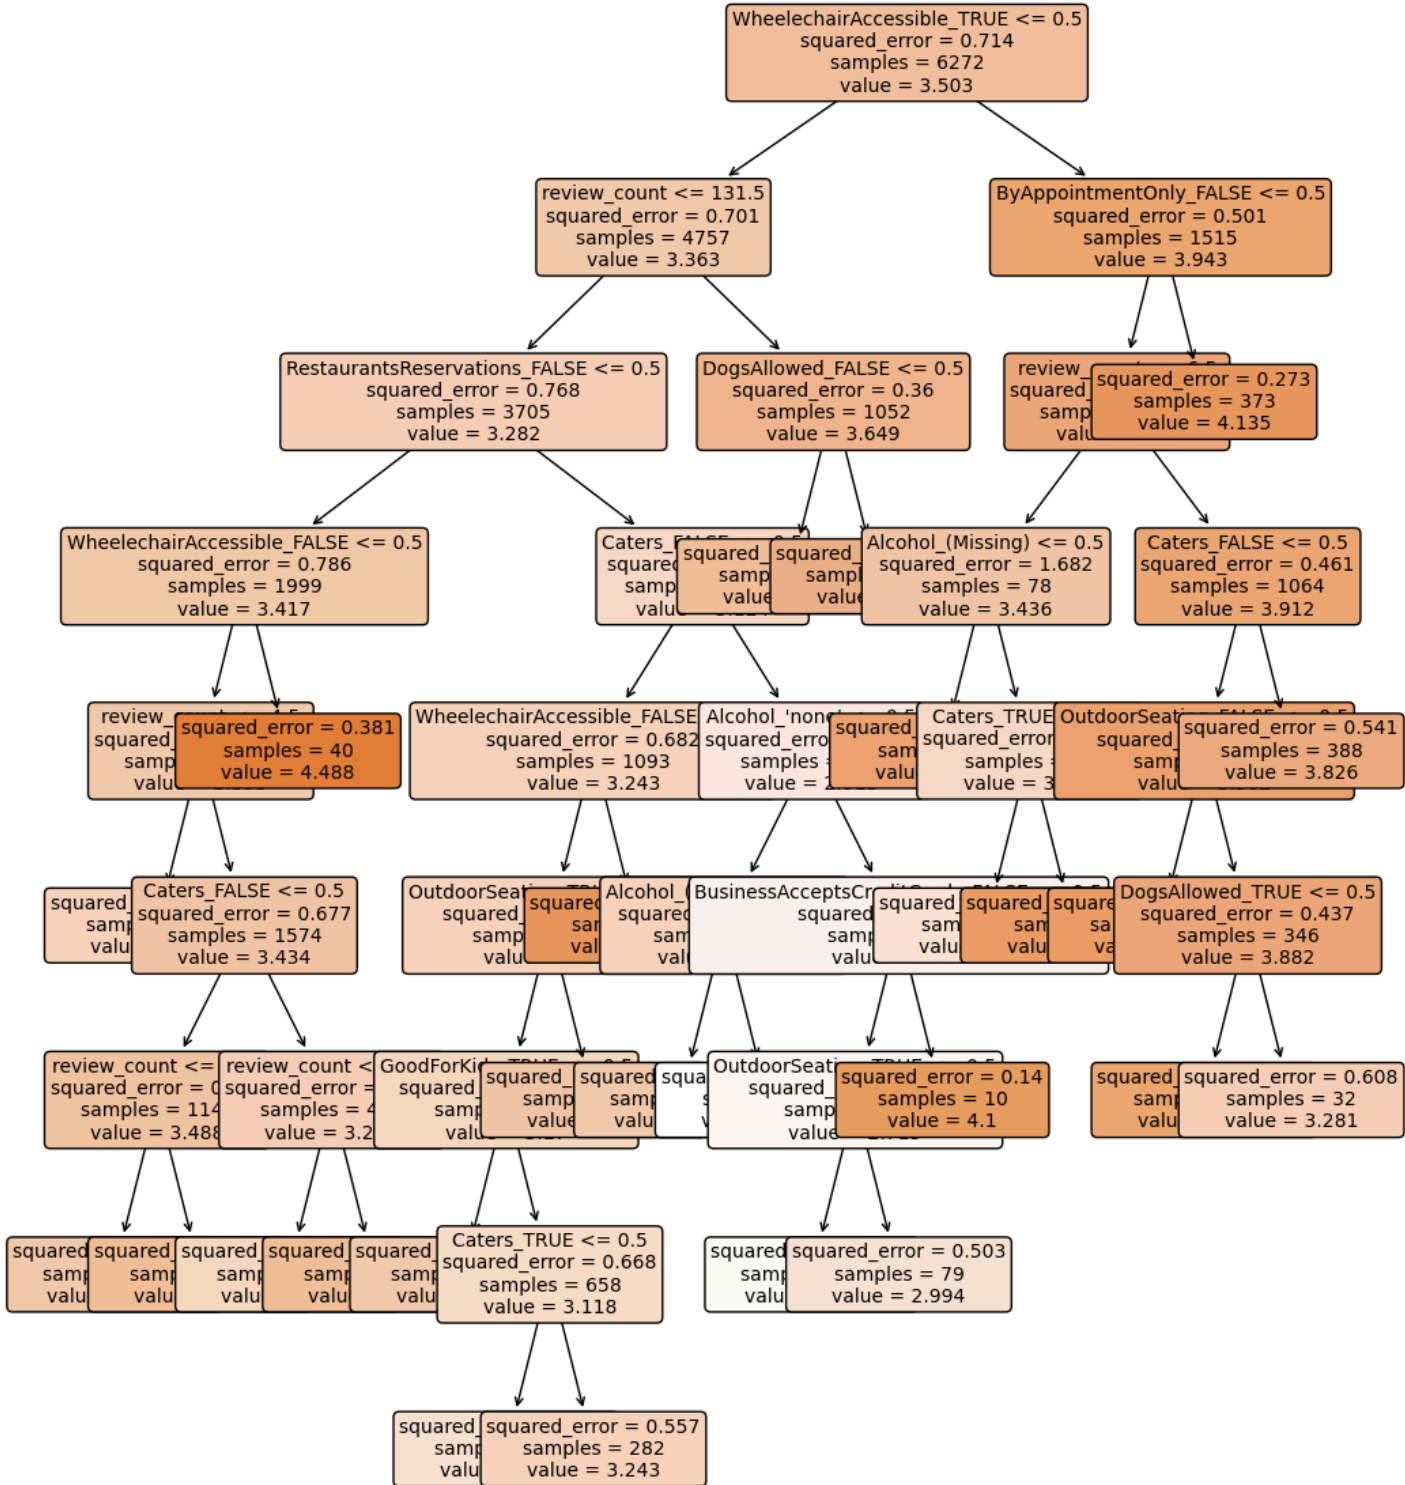




```
max_depth = 15,  
min_samples_split = 25)  
tree = dtrBest.fit(Dummy_train,Ytrain)  
  
#Tree diagram plot  
  
print('Node count =', dtrBest.tree_.node_count)  
plt.figure(figsize=(12,15))  
plot_tree(dtrBest,  
          feature_names = Dummy_train.columns,  
          filled=True,  
          impurity=True,  
          rounded=True,  
          fontsize = 10  
          )  
plt.show()
```

Node count = 51



```
In [9]: #encoding my Test set, and predicting target variables.
        DummyTest = pd.get_dummies(Xtest, columns = [col for col in Xtest.columns if col != "review_count"])

        Ypred = tree.predict(DummyTest)
```

```

In [10]: #functions for OSR^2 and MAE
def OSR2(Ypreds, trueTest):
    RSS = np.sum((trueTest - Ypreds)**2 )
    TSS = np.sum((trueTest - np.mean(trueTest))**2)
    return 1 - (RSS/TSS)
def MAE(Ypreds, trueTest):
    return np.sum(np.abs(Ypreds - trueTest))/len(Ypreds)

In [11]: #the following cells are outputs of OSR2 and MAE of decision tree and regression predictions.

MAE(Ypred, Ytest)

Out[11]: 0.6252132561993073

In [12]: OSR2(Ypred, Ytest)

Out[12]: 0.15099276886806745

In [13]: regrPred = model.predict(Xtest)

In [14]: MAE(regrPred, Ytest)

Out[14]: 0.630129081526346

In [15]: OSR2(regrPred, Ytest)

Out[15]: 0.1625081247047695

In [16]: #make copies so we can establish a new target variable,boolean of above/ below 4 star rating.
trainGT4, testGT4 = train.copy(), test.copy()

In [17]: #make bool columns if a row has a rating of >= 4
trainGT4["fourOrAbove"] = (trainGT4["stars"] >= 4).astype(int)
testGT4["fourOrAbove"] = (testGT4["stars"] >=4).astype(int)

In [18]: #2d ii:
# Code for converting predictions of at least four/ below 4 into boolean outcomes

def regr4orAbove(x):
    x = np.asarray(x)
    return [1 if i >= 4 else 0 for i in x ]

def logitYhat(x):
    x = np.asarray(x)
    return [1 if i>= 0.35272055398 else 0 for i in x]

In [19]: #Logistic regression model for Four or above rating classifier

logistic40A = smf.logit('fourOrAbove ~ review_count + C(GoodForKids, Treatment(reference="(Missing)"
+ C(BusinessAcceptsCreditCards, Treatment(reference="(Missing)")) + C(WiFi, Treat
+ C(BikeParking, Treatment(reference="(Missing)")) + C(ByAppointmentOnly, Treatme
+ C(WheechairAccessible, Treatment(reference="(Missing)")) + C(OutdoorSeating,
+ C(RestaurantsReservations, Treatment(reference="(Missing)")) + C(DogsAllowed,
+ C(Caters, Treatment(reference="(Missing)"))', data = trainGT4).fit()

```

Optimization terminated successfully.  
Current function value: 0.607102  
Iterations 6

```
In [20]: #predictions of the logistic regression function.
logitPred = logistic40A.predict(Xtest)
logitPred = 1/(1+ np.e**logitPred)
logitPred
```

```
Out[20]: 0      0.414209
1      0.317227
2      0.308660
3      0.354872
4      0.422787
...
2683   0.396007
2684   0.437190
2685   0.319350
2686   0.329078
2687   0.376931
Length: 2688, dtype: float64
```

```
In [21]: #dataframe cleaning, etc.
trainGT4X = trainGT4.drop(columns = ["stars","fourOrAbove"])
testGT4X = testGT4.drop(columns = ["stars","fourOrAbove"])
testGT4X = pd.get_dummies(testGT4X, columns = [col for col in Xtest.columns if col != "review_co
testGT4Y = testGT4["fourOrAbove"]
trainGT4Y = trainGT4["fourOrAbove"]
trainGT4X = pd.get_dummies(trainGT4X, columns = [col for col in Xtest.columns if col!= "review_c
```

```
In [22]: #Cross validation and model fitting for classification tree
grid_values = {'ccp_alpha': np.linspace(0,0.01,26),
               'min_samples_leaf': [3,4,5,6,7],
               'min_samples_split': [15,20,25],
               'max_depth': np.arange(12,20),
               'class_weight' : ["balanced"],
               }

dtc = DecisionTreeClassifier()
dtcCV = GridSearchCV(dtc, param_grid = grid_values, scoring = 'accuracy', cv=4, verbose=1,n_jobs
dtcCV.fit(trainGT4X, trainGT4Y)
```

Fitting 4 folds for each of 3120 candidates, totalling 12480 fits

```
Out[22]: ▸ GridSearchCV
▸ estimator: DecisionTreeClassifier
    ▸ DecisionTreeClassifier
```

```
In [23]: #tree diagram for classifier
print(dtcCV.best_params_)
print('Node count =', dtcCV.best_estimator_.tree_.node_count)
plt.figure(figsize=(12,18))
plot_tree(dtrBest,
          feature_names = trainGT4X.columns,
          filled=True,
          impurity=True,
          rounded=True,
          fontsize = 10
```

```
)  
plt.show()
```

```
{'ccp_alpha': 0.0008, 'class_weight': 'balanced', 'max_depth': 12, 'min_samples_leaf': 3, 'min_samples_split': 15}  
Node count = 45
```



```
Out[24]: {'ccp_alpha': 0.0008,
          'class_weight': 'balanced',
          'max_depth': 12,
          'min_samples_leaf': 3,
          'min_samples_split': 15}
```

```
In [25]: #predict the Above 4 star rating for each establishment using the classification tree.
Ypred2 = dtcCV.best_estimator_.predict(testGT4X)
Ypred2
```

```
Out[25]: array([0, 1, 1, ..., 1, 1, 1])
```

```
In [26]: #baseline naive classifier, predicts mode of the outcomes.
y_baseline = testGT4Y.mode()
modePred = pd.Series([y_baseline[0] for _ in range(len(testGT4Y))])
```

```
In [27]: #function that returns a 3-list with accuracy, TPR, and FPR given a true y values array and a predicted y values array
def AccTFR(y,yhat):
    y = np.asarray(y)
    yhat = np.asarray(yhat)
    return [np.mean(y == yhat), np.sum((y == 1) & (yhat == 1)) / (np.sum((y == 1) & (yhat == 1)))]
```

```
In [28]: # constructing the dataframe "table" for each of the five models, with accuracy, tpr, and fpr being the rows and the models as the columns

df = pd.DataFrame()

df["baseline"] = AccTFR(testGT4Y,modePred)
df["lin. regression Threshold"] = AccTFR(testGT4Y, regr4orAbove(r2d2))
df["regressor Dec. Tree"] = AccTFR(testGT4Y, regr4orAbove(Ypred))
df["log. Regression"] = AccTFR(testGT4Y, logitYhat(logitPred))
df["Classification Tree"] = AccTFR(testGT4Y, Ypred2)

df.index = ["Accuracy:", "TPR:", "FPR:"]
df
```

```
Out[28]:
```

	baseline	lin. regression Threshold	regressor Dec. Tree	log. Regression	Classification Tree
<b>Accuracy:</b>	0.551339	0.616071	0.616443	0.338542	0.643229
<b>TPR:</b>	0.000000	0.203980	0.202322	0.628524	0.538143
<b>FPR:</b>	0.000000	0.048583	0.046559	0.897436	0.271255

```
In [29]: #question 2e- feature importance
imp = dtcCV.best_estimator_.feature_importances_
features = Dummy_train.columns
importance_df = pd.DataFrame({
    'Feature': features,
    'Importance': imp
})

# Sort the DataFrame by importance
importance_df = importance_df.sort_values(by='Importance', ascending=False)

importance_df
```



Out[29]:

	Feature	Importance
16	WheeelchairAccessible_TRUE	0.488146
0	review_count	0.195819
15	WheeelchairAccessible_FALSE	0.068601
19	RestaurantsReservations_FALSE	0.065059
23	Caters_FALSE	0.052794
3	Alcohol_'full_bar'	0.024468
13	ByAppointmentOnly_FALSE	0.022670
18	OutdoorSeating_TRUE	0.022189
6	BusinessAcceptsCreditCards_FALSE	0.021529
21	DogsAllowed_FALSE	0.014867
2	GoodForKids_TRUE	0.013512
4	Alcohol_'none'	0.010345
22	DogsAllowed_TRUE	0.000000
20	RestaurantsReservations_TRUE	0.000000
17	OutdoorSeating_FALSE	0.000000
12	BikeParking_TRUE	0.000000
14	ByAppointmentOnly_TRUE	0.000000
1	GoodForKids_FALSE	0.000000
11	BikeParking_FALSE	0.000000
10	WiFi_(Missing)	0.000000
9	WiFi_'paid'	0.000000
8	WiFi_'no'	0.000000
7	BusinessAcceptsCreditCards_TRUE	0.000000
5	Alcohol_(Missing)	0.000000
24	Caters_TRUE	0.000000